

## Why GFlowNet?

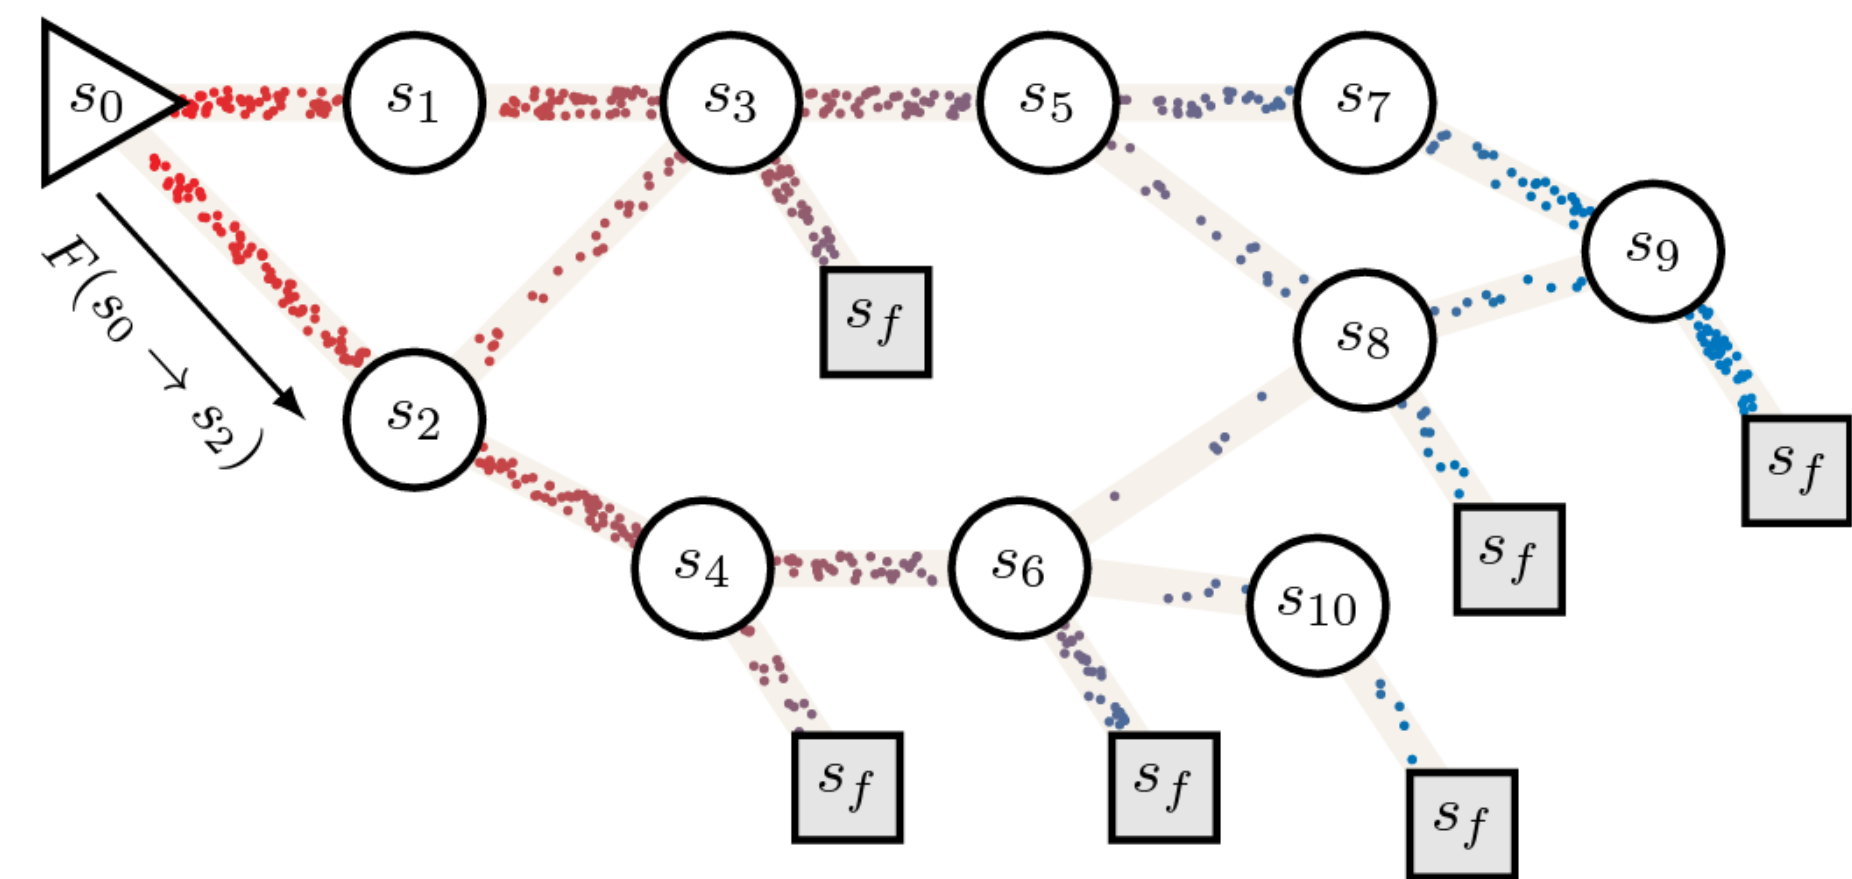
GFlowNets learn a generative policy that produces samples in a single forward pass, avoiding the limitations of long Markov chains in MCMC. Key advantages include:

- **Amortized Sampling:** Once trained, GFlowNets generate independent samples efficiently, without sequential mixing.
- **Diversity:** GFlowNets explore multiple modes of the solution space, producing diverse samples proportional to the reward function  $R(s)$ .
- **Scalability:** By leveraging neural networks for function approximation, GFlowNets can efficiently scale to large state spaces.

These properties are well suited for our redistricting problem.

## GFlowNet Overview

GFlowNets generate structured objects (e.g., redistricting plans) step-by-step by sampling actions through a Directed Acyclic Graph (DAG):



- **States and Actions:** Starting from an initial state  $s_0$ , actions transition the system to new states, ending in a final state where the object is complete.
- **Flow Conservation:** The sum of flows into a state equals the sum of flows out, ensuring terminal state probabilities are proportional to their rewards  $R(s_f)$ .
- **Learning Objective:** GFlowNets optimize a generative policy to match the sampling distribution to the reward function by minimizing a flow-matching loss.
- **Diversity:** Multiple trajectories to terminal states enable exploration of diverse, high-reward solutions.

## Design Choices for Our GFlowNet Implementation

**State Representation:** Each state  $\mathcal{S}$  is represented as a tensor of shape:

$$\mathcal{S} \in \mathbb{N}^{N_{\text{counties}} \times (L_{\text{adj}} + 1 + \text{num\_district})},$$

where  $N_{\text{counties}}$  is the number of counties, and  $L_{\text{adj}}$  is the maximum adjacency length.

- Each row corresponds to a county.
- The first  $L_{\text{adj}}$  entries encode the adjacency information, padded to a fixed length.
- The entry  $L_{\text{adj}} + 1$  stores the current district assignment.
- The last  $\text{num\_district}$  entries are to encode valid actions to help the model learn faster.

Counties located at district borders have their district IDs marked negative to highlight valid actions for the MLP.

**Action Representation:** Actions  $\mathcal{A}(s)$  at state  $s$  are represented as a vector of shape:

$$\mathcal{A}(s) \in [V_{\text{ID}}, N_{\text{ID}}] \in \mathbb{N}^2.$$

Here,  $V_{\text{ID}}$  is the county identifier, and  $N_{\text{ID}}$  is the target district for reassignment.

**Ensuring DAG Structure:** To enforce acyclic state transitions, we enforce that a county cannot return to a previously assigned district. This ensures that redistricting plans evolve incrementally along a Directed Acyclic Graph (DAG), with each transition forming a valid step toward the final configuration.

**Initial State  $s_0$ :**

The initial state  $s_0$  is the 2020 district partition. Ensuring that all counties are initially assigned to valid districts while preserving adjacency constraints.

## Reward Definition

Our reward function is a weighted combination of previously shown evaluation metrics with tunable coefficients  $\lambda$ :

$$\begin{aligned} R_{s_f} &= \lambda_1 \cdot \text{Entropy\_pop} \\ &+ \lambda_2 \cdot \text{Compactness} \\ &+ \lambda_3 \cdot (1 - |\text{Partisan Bias}|) \\ &+ \lambda_4 \cdot (1 - |\text{Eff\_Gap}|) \end{aligned}$$

Invalid actions are heavily penalized to guide the MLP toward learning valid transitions.

## MCMC Results

### Partisan Bias vs % Changed Precincts

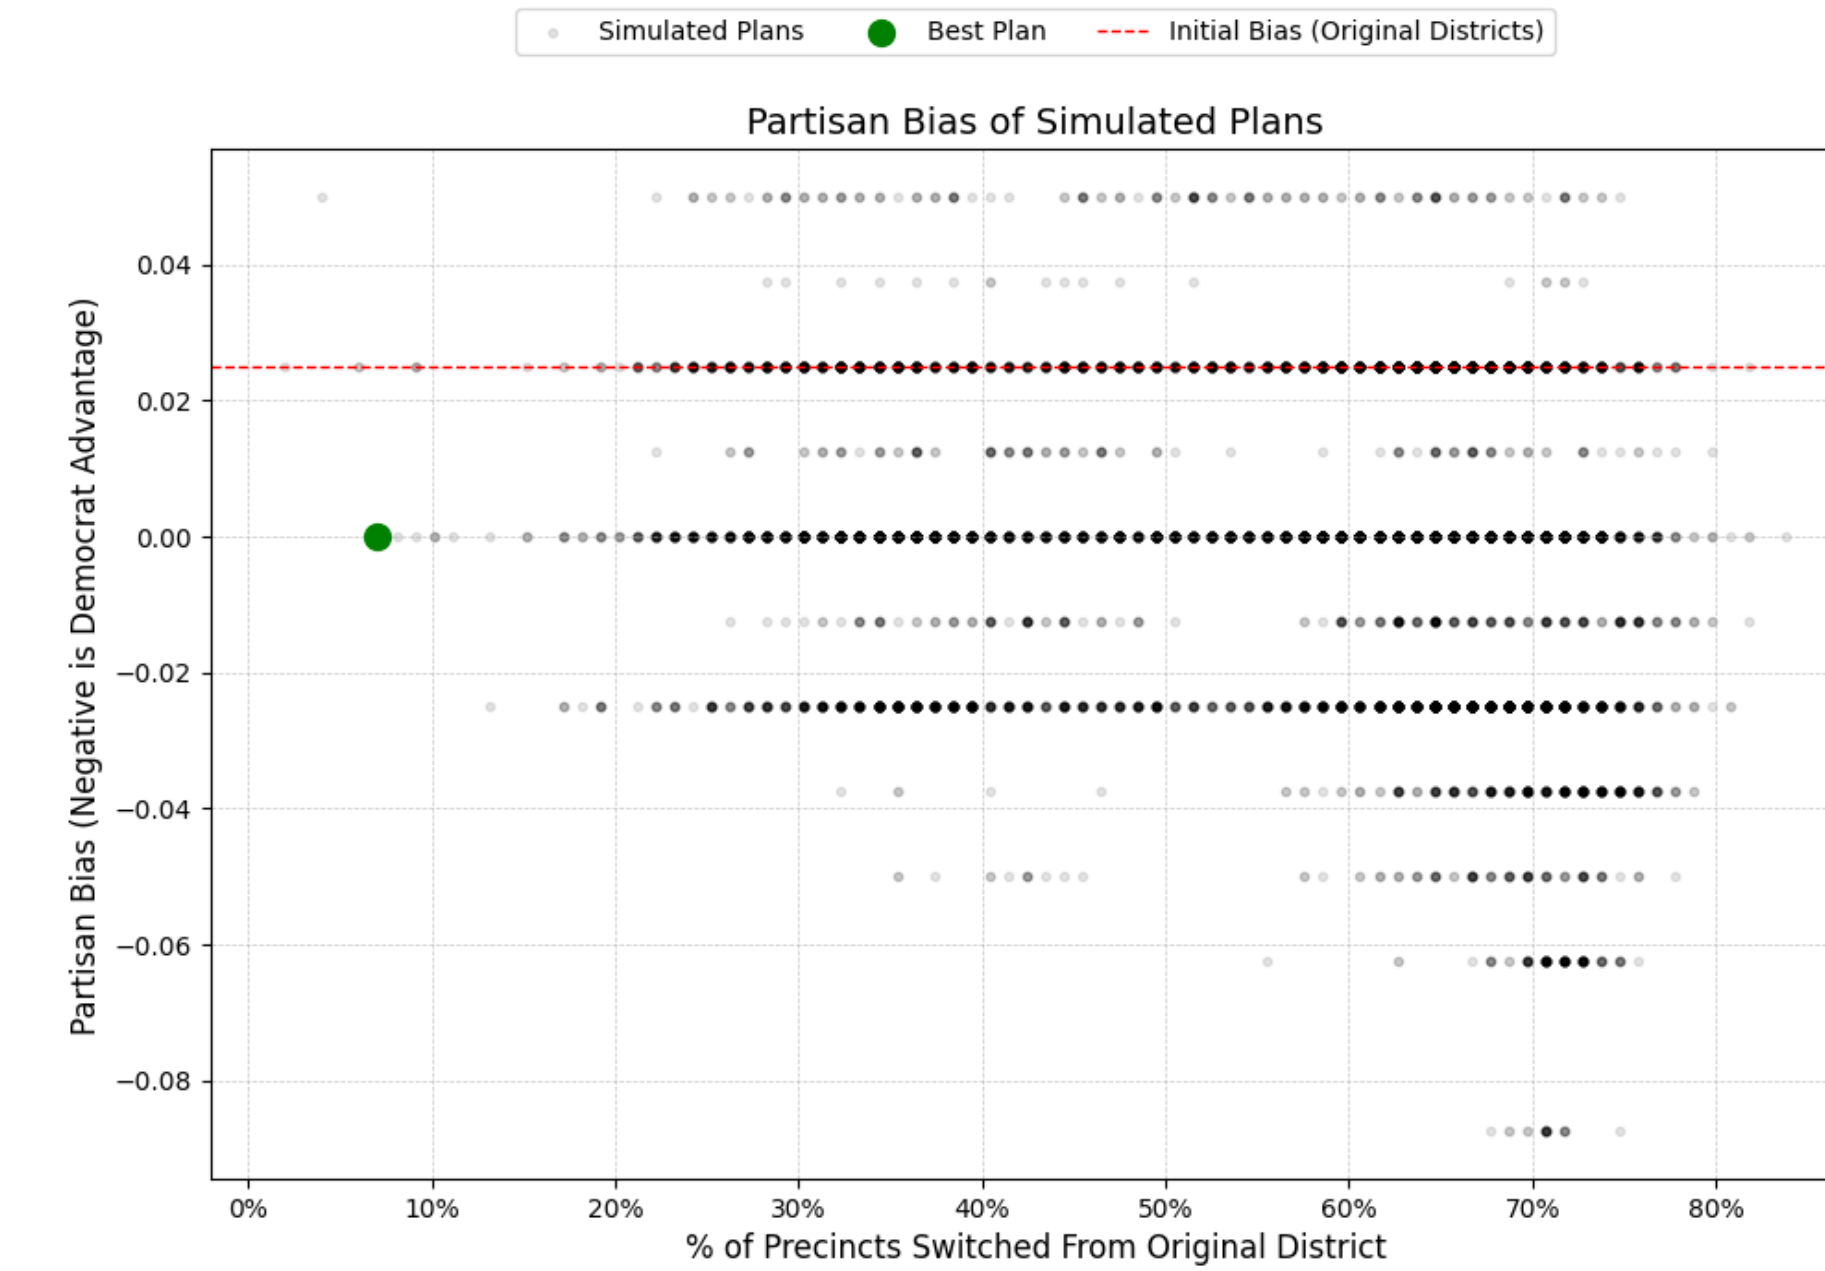


Figure 1: Partisan Bias Across Precinct Swaps

The best plan achieved zero partisan bias (green dot) while the original districts maintained a bias of 0.0250 (red dashed line). while enforcing a minimal change of only 5% of the original presincts.

### Map Comparison: MCMC

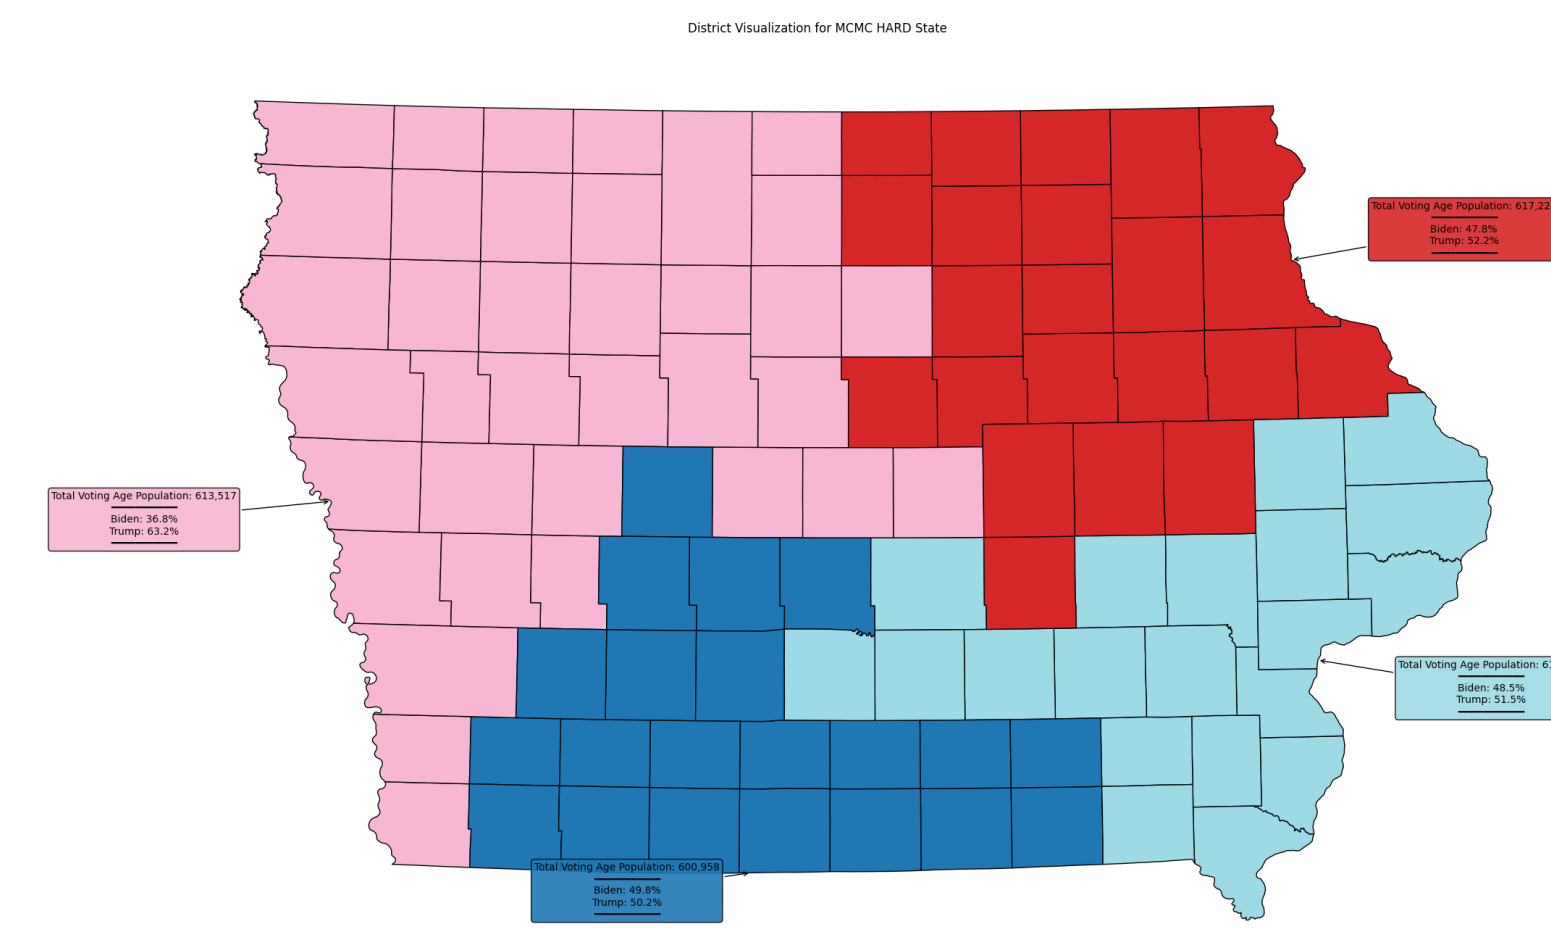


Figure 2: Original District Map (2020)

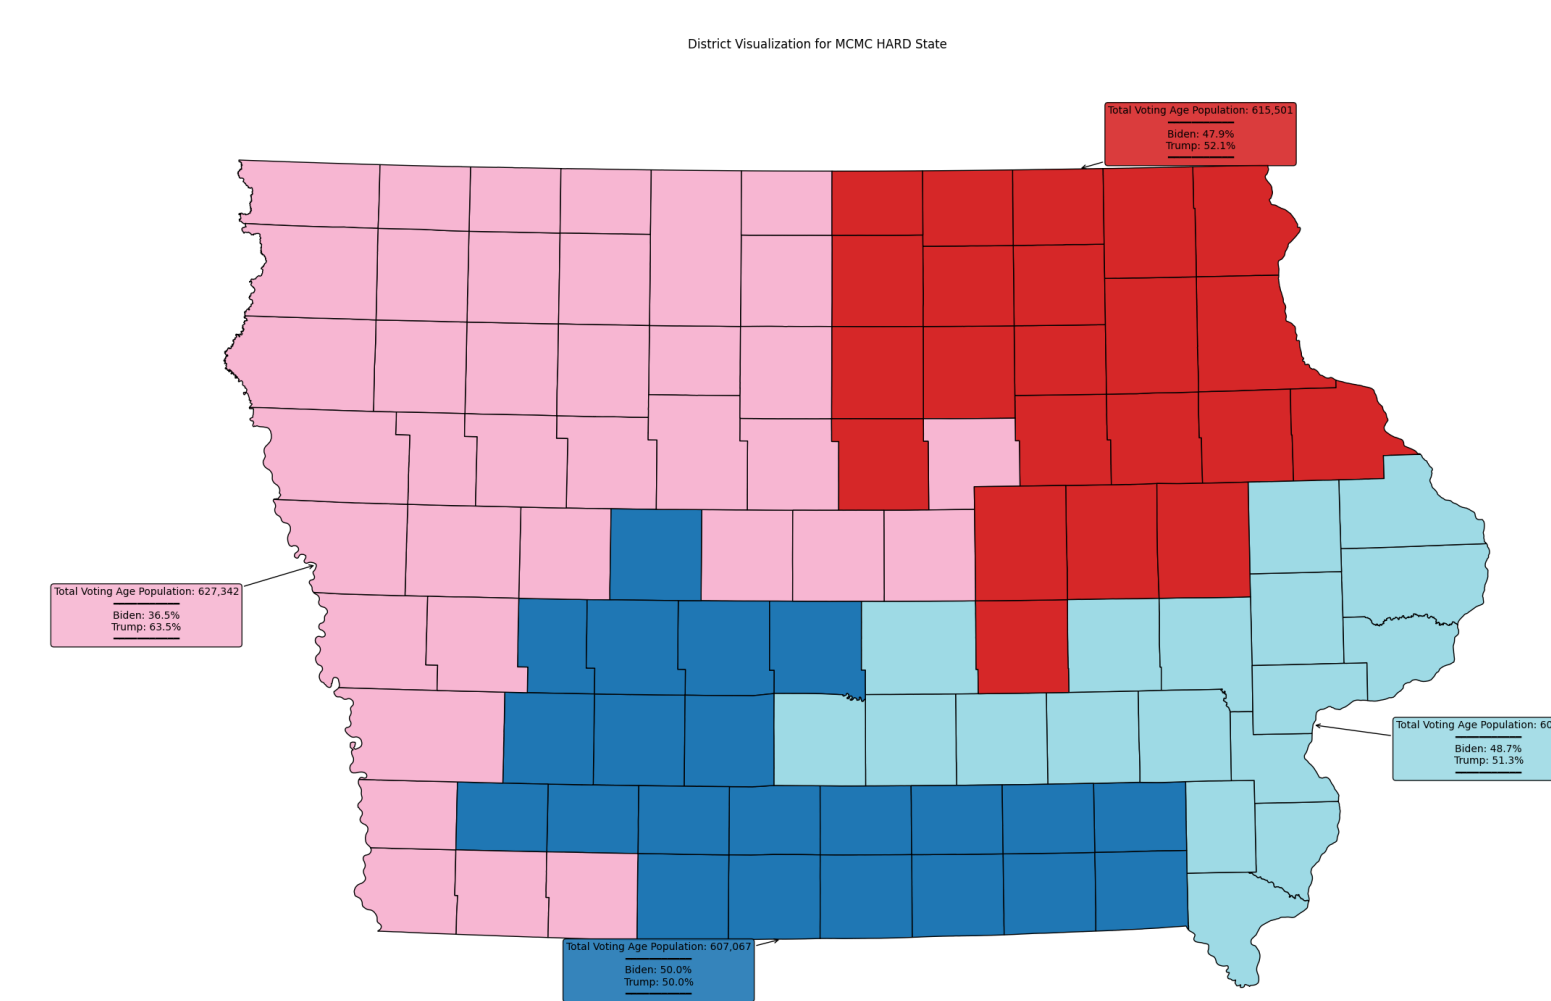


Figure 3: Optimized District Map Using MCMC

### Comparison of Metrics:

Metric	Optimized	Original	GFN
Partisan Bias	<b>0.0000</b>	0.0250	<b>0.0000</b>
Efficiency Gap	0.1641	0.4163	<b>0.0909</b>
Mean Compactness	0.2962	0.3334	0.2851
Compactness Std	0.0611	0.0574	0.0555
Max Pop. Deviation	2.8679%	0.0066%	3.535%

Table 1: Comparison of Districting Metrics

## GFN Results

We trained our GFN on 128,000 trajectories. The policy samples actions until it picks an invalid one or the exit action. Penalties are given for invalid actions and our reward metrics is computed at each state. We also introduced some probability  $\epsilon = 0.03$  of the sampled action to be taken from a uniform distribution over actions.

### Map Comparison: GFN

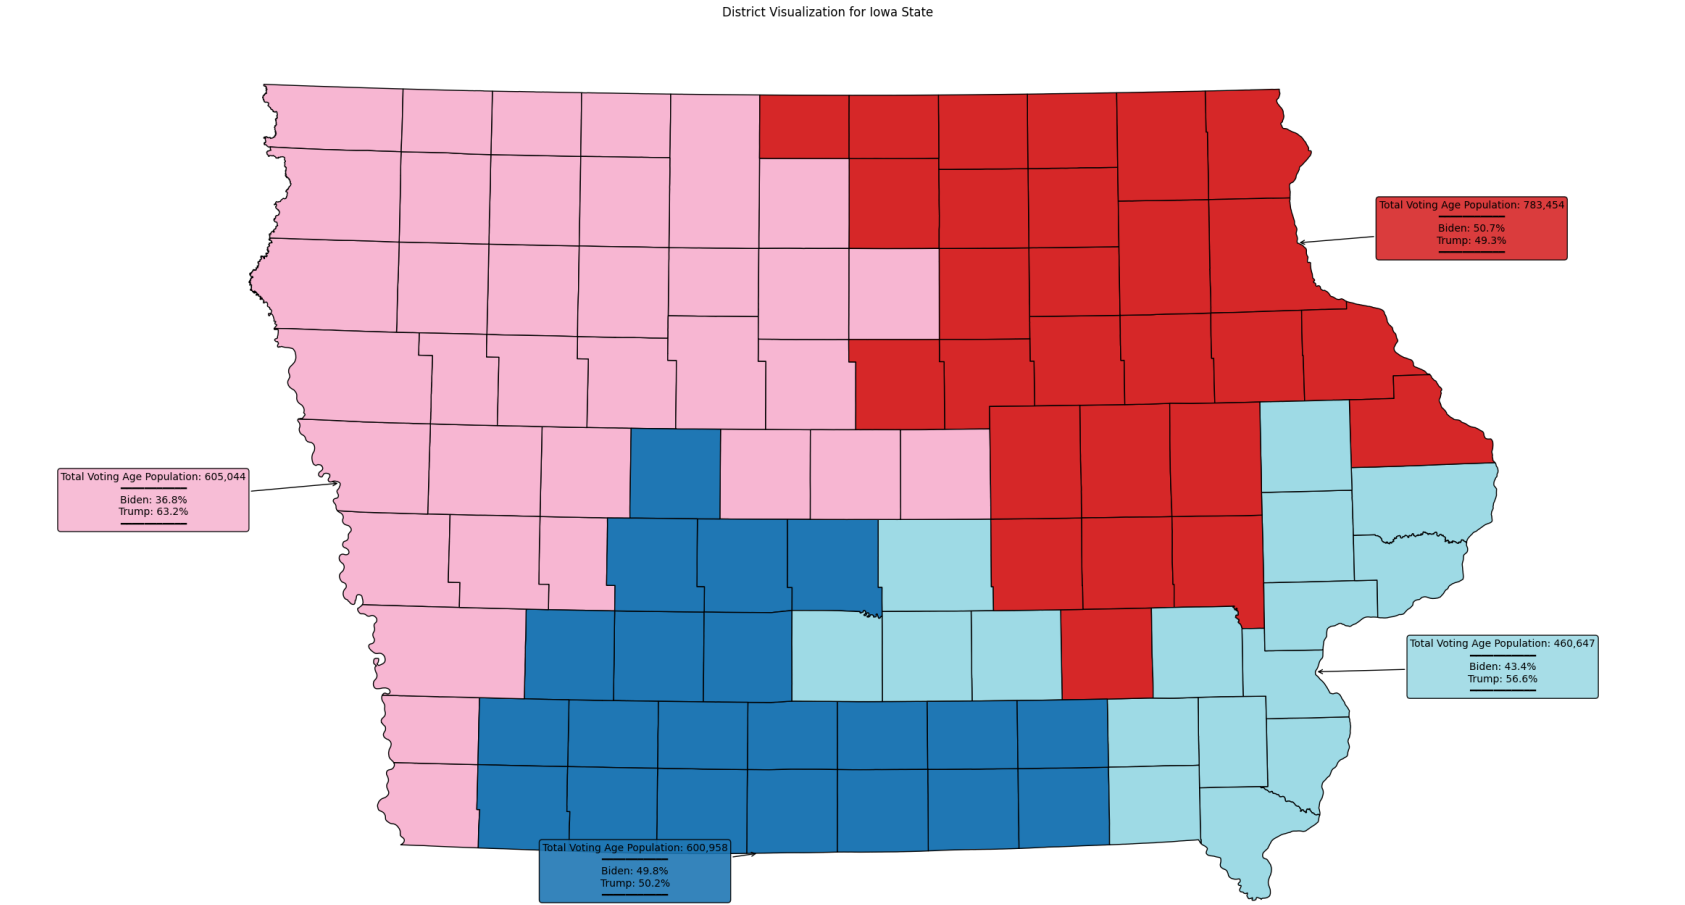


Figure 4: Optimized District Map Using GFN

Our GFlowNet method yields better results than the MCMC and amortizes training so that training the model on an other state would not require to sample as many trajectories to achieve optimal performance.

### Possible Improvements

We believe the generalization capability of the GFN could be improved by considering random starting states for each trajectories.

### Conclusion

Our results show the advantages of using GFlowNets for graph exploration in the context of Gerrymandering. We demonstrate more efficient exploration that yields to fairer district partitions of electoral counties. These methods could be use to ensure fairness in repartitioning often used to shift political advantage.